

# Design a Notification System

## 1. Main Types of Notifications

### 1. Push Notifications

- Sent to mobile apps / browsers
- Example:
  - “Your order has been shipped”
  - “New message received”

### 2. Email Notifications

- Sent through email services (SMTP, SES, SendGrid)
- Example:
  - Account verification
  - Monthly reports
  - Password reset

### 3. SMS Notifications

- Sent to phone numbers via SMS gateways
- Example:
  - OTP verification
  - Banking alerts

### 4. In-App Notifications

- Shown inside the application UI
- Example:
  - Instagram like notifications
  - LinkedIn connection request

### 5. Webhook / System Notifications (service-to-service)

- One system notifying another system
- Example:
  - Payment success webhook
  - Order completed event

## Another Important Classification

### Based on timing

1. Real-time notifications
  - OTP, chat message
2. Scheduled notifications

- Daily reminders
3. Batch notifications
- Weekly reports sent to millions of users

When interviewer says “Design a Notification System”, you should ask these questions.

## 1. Product Requirement Questions

- What types of notifications should we support?
  - Email, SMS, Push, In-App?
- Is the system real-time, scheduled, or batch notifications?
- Should users be able to opt-in / opt-out of notifications?
- Do we need user preference settings (mute categories, quiet hours)?

## 2. Scale Questions

- How many daily active users?
- How many notifications per second expected?
- What is peak traffic?

## 3. Reliability Questions

- Should notifications be guaranteed delivery or best effort?
- Do we need retry mechanisms if sending fails?
- Should notifications be deduplicated (avoid sending twice)?

## 4. Latency Questions

- What is acceptable delivery delay?
  - milliseconds?
  - seconds?
  - minutes?

## 5. Storage Questions

- Do we need to store notification history?
- For how long should notifications be retained?

## 6. Third-party integration

- Are we using external providers (Firebase, SES, Twilio)?
- Should system support multiple providers fallback?

# API Interactions

## 1. Retrieve User Info

Used by notification service to get delivery details (email, phone, device token).

API

GET /users/{user\_id}

Response

```
{  
  "user_id": "123",  
  "email": "user@gmail.com",  
  "phone": "+919876543210",  
  "device_token": "abcxyz"  
}
```

## 2. Send Notification

Used by other services (order service, payment service, etc.) to trigger notifications.

API

POST /notifications

Request

```
{  
  "user_id": "123",  
  "type": "push",  
  "title": "Order Shipped",  
  "message": "Your order #123 has been shipped"  
}
```

Response

```
{  
  "notification_id": "n789",  
  ...  
}
```

```
"status": "queued"
```

```
}
```

### 3. Fetch Notification History

User fetches all past notifications.

API

```
GET /users/{user_id}/notifications
```

Optional pagination:

```
GET /users/{user_id}/notifications?page=1&limit=20
```

Response

```
[
```

```
{
```

```
  "notification_id": "n1",
```

```
  "title": "Payment Successful",
```

```
  "message": "Payment of ₹500 completed",
```

```
  "status": "delivered",
```

```
  "timestamp": "2026-02-09T10:00:00Z"
```

```
}
```

```
]
```

These are called third-party notification delivery providers.

They are external services that actually deliver the notification to the user's device.

Your notification system does not directly send SMS/email/push – instead it calls these providers' APIs.

## Examples

### 1. Push Notification Providers

- APNs (Apple Push Notification service) → sends notifications to iOS devices
- FCM (Firebase Cloud Messaging) → sends notifications to Android / Web

### 2. SMS Providers

- Twilio
- Nexmo (Vonage)

Used to send:

- OTP messages
- alerts
- transactional SMS

### 3. Email Providers

- Mailchimp
- SendGrid
- Amazon SES

Used for:

- account emails
- marketing emails
- transactional emails

## How it works in architecture

Application Service




Notification Service



Third-party provider (FCM / APNs / Twilio / SES)



User device



It represents a centralized Notification Server that receives notification requests from multiple internal services and sends them to users through different delivery providers.

## Flow Explanation

### 1. Producer Services

- Service 1, Service 2, Service 3
- These are application microservices (Order service, Payment service, Social service, etc.)
- They call the Notification Server API to request sending a notification.

### 2. Notification Server

- Central system responsible for:
  - validating requests
  - selecting notification channel (push, SMS, email)
  - formatting message
  - forwarding to correct provider

### 3. Third-Party Providers

- APNs → iOS push notifications
- FCM → Android / Web push notifications
- Twilio → SMS
- Mailchimp → Email

These providers actually deliver the notification to the client device.

### 4. Cron Jobs / Scheduled APIs

- Used for scheduled notifications
  - reminders
  - daily summaries
  - marketing campaigns

## Problems in this Design

### 1. No Queue (Biggest issue)

Requests go directly from services → Notification Server → providers.

Problems:

- Traffic spikes can overload notification server
- Requests may fail if provider is slow
- No buffering mechanism

## 2. Single Notification Server = Single Point of Failure

If the notification server crashes:

- All notifications stop
- Entire system becomes unavailable

## 3. No Retry Mechanism

If:

- Twilio fails
- FCM temporarily down

Notifications will be lost because no retry logic exists.

## 4. Tight Coupling with Providers

Notification server directly calling:

- APNs
- FCM
- Twilio
- Mailchimp

Any provider latency will block processing.


## 5. No Scalability


As traffic increases:

- One server cannot handle millions of notifications/sec
- Horizontal scaling difficult without queues/workers

## 6. No Prioritization / Scheduling Handling

- Urgent OTP and marketing messages treated the same
- Scheduled notifications not clearly handled





## Step-by-step Explanation

### 1. Producer Services (Left side)

- Service 1, Service 2 ... Service N
- These services generate notification events (order placed, payment success, message received).
- They send requests to the Notification Publisher Service.

### 2. Publisher (Notification Service)

Responsibilities:

- Validate request
- Determine notification channel (Push / SMS / Email)
- Store notification in Notification DB
- Use Cache for fast user/device lookup
- Publish message into appropriate Message Queue (MQ)

Example:

- Push notification → Push MQ
- SMS → SMS MQ
- Email → Email MQ

### 3. Message Queues (MQ)

- Act as buffer
- Handle traffic spikes
- Ensure notifications are not lost
- Enable asynchronous processing

- Allow horizontal scaling

Each channel has its own queue.

#### 4. Workers (Consumers)

- Workers read messages from MQ
- Format and send notification to the appropriate provider
  - APNs → iOS
  - FCM → Android/Web
  - Twilio → SMS
  - Mailchimp → Email
- Multiple workers can run in parallel for scaling.

#### 5. Delivery Providers → Client

Providers deliver the notification to the user device.

### Why this design is better



- No single bottleneck
- Handles traffic spikes
- Supports retries
- Horizontally scalable
- Reliable delivery
- Independent scaling per channel (email, sms, push)

Cache & DB : user Info, Notification templates etc

MQ : Removed dependency bw 3rd party APIs and our notification server. Service is now Async

Improvements:

1. Notification Templates
2. Notification settings on user level :  
user can turn on/off a particular type of notification
3. Rate Limiting
4. API authentication
5. Monitoring Logic
6. Notification Log.



# 1. Producer Services

- Service 1, Service 2 ... Service N
- These services generate events:
  - order placed
  - payment success
  - new message
- They call the Notification Publisher API

# 2. Publisher / Notification Service

This is the entry point of the system.

Responsibilities:

- Authentication & Authorization
- Rate limiting (protect system from abuse)
- Validate notification request
- Determine channel (push / email / sms)
- Store notification in Notification DB

- Use Cache for fast user/device lookup
- Push message to appropriate Message Queue (MQ)

## 3. Channel-specific Message Queues (MQ)

Separate queues for:

- Push notifications
- SMS
- Email

Why:

- Independent scaling
- Isolation of failures
- Prioritization possible

Queues act as buffer during traffic spikes.

## 4. Workers (Consumers)

Workers read messages from MQ:

- Format message
- Call third-party providers:
  - APNs → iOS push
  - FCM → Android push
  - Twilio → SMS
  - Mailchimp → Email
- Update Notification Log with delivery status

Multiple workers run in parallel for scaling.

## 5. Notification Log

Stores:

- delivery status
- retry attempts
- timestamps

Used for:

- debugging
- user notification history
- retry handling

## **6. Analytics Service**

**Collects data:**

- notifications sent
- open rates
- click rates
- failures
- campaign metrics

**Helps business dashboards and reporting.**

## **7. End Delivery**

**Third-party providers finally deliver notifications to client devices.**